

# MATH 5490 — HPC-xTC

Home Work #2

Group 5 JORGE FLORES, XIUKUN HU, GRIGORII SARNITSKII

## 1 Step 2

### 1.1 Results

The code is capable with square blocks (i.e. with `blk_rows` and `blk_cols` to be the same) and multiple blocks in A and B. The multiplication result will be written to disk in file(s) `C.*.*`, and can be printed on screen.

The result is examed by MATLAB for  $100 \times 100$ ,  $1000 \times 1000$  and  $5000 \times 5000$  block multiplication. The absolute difference between two results is less than  $1e - 11$  for all tests.

Total time spent for  $5000 \times 5000$  matrices multiplication for this code is  $1.37e3$  seconds and computing time is  $1.36e3$  seconds, while for MATLAB the computing time is only 4.28 seconds.

### 1.2 Requirement

Unzip the folder and every file needed is inside. Use `make run` command to run it.

### 1.3 Compute Bound

For MacBook Pro with 2.7GHz Intel Core i5, 8G DDR3 memory, it gets compute bound when the block is over  $50 \times 50$ .

For Lenovo with 2.4GHz Intel Core i7, running in a 64bit Red Hat virtual machine, it gets compute bound when the block is over  $50 \times 50$ .

## 2 Step 3

### 2.1 Algorithm

The pseudocode is as follow:

## OMP PARALLEL

### OMP single

$ablock[tog] \leftarrow READ(A\_0\_0)$

▷ tog initialized as 0

$bblock[tog] \leftarrow READ(B\_0\_0)$

### ENDOMP single

▷ Implicit barrier here

**while**  $i < \text{rows of blocks in } C$  **do**

### OMP single nowait

$ablock[1 - tog] \leftarrow \text{READ}(\text{next } A \text{ block})$

$bblock[1 - tog] \leftarrow \text{READ}(\text{next } B \text{ block})$

### ENDOMP single nowait

### OMP for nowait dynamic

**for** every continuous  $WIDTH$  (macro) elements in  $ablock$  **do**

**for** each column of  $B$  **do**

$temp+ = WIDTH$  elements of  $ablock[tog]*WIDTH$  rows of  $bblock[tog]$

### OMP atomic update

$cblock[ctog]$  corresponding element  $+ = temp$

### ENDOMP atomic update

**end for**

**end for**

### ENDOMP for nowait dynamic

**if**  $blk\_cols$  cannot be divided by  $WIDTH$  **then**

### OMP for nowait dynamic

**for** each row of the remainder columns in  $ablock[tog]$  **do**

update  $cblock[ctog]$

**end for**

### ENDOMP for nowait dynamic

**end if**

## OMP BARRIER

**if**  $k == \text{columns of blocks in } A$  **then**

### OMP single nowait

$WRITE\ cblock[ctog]$

fill  $cblock[ctog]$  with zeros

### ENDOMP single nowait

$ctog = 1 - ctog$

**end if**

$tog = 1 - tog$

update  $i, j, k$  to point to next block in  $A$  and  $B$ ;

**end while**

## ENDOMP PARALLEL

## 2.2 Results

The code works for square blocks (i.e. with `blk_rows` and `blk_cols` being equal) and multiple blocks in A and B. The multiplication result will be written to disk in file(s) of the form `C.*.*`.

The result was verified to be correct using MATLAB for  $1000 \times 1000$  block size and  $10 \times 10$  blocks matrix multiplication.

We performed a one dimensional analysis on the performance of our code using different block sizes and block matrix sizes. These tests were performed on an Intel(R) Xeon(R) CPU E5-4650 @2.70GHz with 8 cores and 2 threads per core (for a total of 16 threads). The results are summarized in the table below, where the total time represents the time it took our code to multiply the matrices and write the resulting blocks of  $C$  to disk.

Block/Matrix Size Analysis		
Block Size	Block Matrix Size	Total Time (seconds)
100x100	1x1	$1.130199 \times 10^{-2}$
100x100	2x2	$5.166101 \times 10^{-2}$
100x100	5x5	$6.825109 \times 10^{-1}$
1000x1000	1x1	$2.701371 \times 10^{-1}$
1000x1000	2x2	1.635549
1000x1000	5x5	22.65355

It is obvious that the total time is  $O(n^3)$ .

We also performed a one dimensional analysis, this time focusing on the performance of our code using different thread numbers. These tests were performed on an Intel(R) Xeon(R) CPU E5-4650 @2.70GHz with 8 cores and 2 threads per core (for a total of 16 threads). The results are summarized in the table below.

Thread Number Analysis			
Number of Threads	Block Block Size	Block Matrix Size	Total Time (seconds)
4	1000x1000	5x5	43.33175
8	1000x1000	5x5	22.61952
12	1000x1000	5x5	20.78736
16	1000x1000	5x5	23.37497

We note that increasing the thread number from 4 threads to 8 threads, the total time nearly halved. However, changing from 8 to 12 and 16 threads, we see only a small change in the total time. This might be due to the fact that the system we used only has 8 physical cores.

## 2.3 Requirement

Unzip the folder called hw2.zip, every file needed to run the code is inside. Use `make matrix` to first generate and write to disk the matrices (block entries for each matrix) that will be used. (Notice that the block should be square.) Then use `make run` to multiply the matrices in parallel and write the resulting matrix (in blocks) to the disk. Alternatively, use `make runserial` to multiply the matrices in serial.

## 2.4 Compute Bound

For a MacBook Pro with 2.7GHz Intel Core i5, 8G DDR3 memory, compute bound is achieved when the block size is over  $40 \times 40$ .

For a Lenovo with 2.4GHz Intel Core i7, running in a 64bit Red Hat virtual machine, compute bound is achieved when the block size is over  $40 \times 40$ .

For an Intel(R) Core(TM) i3-6100U CPU @ 2.30GHz running Ubuntu, compute bound was achieved when the block size is over  $50 \times 50$ .

The information is summarized in the following table:

Compute Bound Analysis	
System	Block Size
Intel(R) Core(TM) i7-4700MQ CPU @ 2.40Ghz GNU bash, version 4.2.46(1)-release (x86_64-redhat-linux-gnu), 4 threads	above 40 x 40
Intel(R) Core(TM) i3-6100U CPU @ 2.30GHz gcc version 5.4.0 20160609 (Ubuntu 5.4.0-6ubuntu1 16.04.2)	above 50 x 50
2.7 GHz Intel Core i5 gcc-6 (Homebrew gcc 6.2.0) 6.2.0, OS X Sierra, 4 threads	above 40 x 40